

Project Approach: IO Digital

Alex Hendrickx
Student Bachelor in de Toegepaste Informatica – Applicatieontwikkeling

Project Proposal and Methodological Justification

This document constitutes the Project Approach for the bachelor's thesis project, titled: "A Scalable GA4 Data Pipeline for Portfolio Analysis". It describes the problem statement, objectives, methodology, and planning underlying the realization of an automated data extraction and enrichment platform for iO.

Management Summary

The project addresses the strategic need to unlock the wealth of data within the Google Analytics 4 (GA4) ecosystem of iO's client portfolio – comprising hundreds of accounts and thousands of digital properties. The current, fragmented access to this data hinders cross-client analyses and the development of data-driven insights.

The proposed solution is the realization of a robust and scalable data pipeline. This system automates the entire ETL process:

1. Extract: Systematically retrieving event data from all relevant GA4 properties.
2. Transform: Enriching the raw data with over 14 dimensions, including an AI-driven classification into 25 sectors, geographical data, and user behavior.
3. Load: Persisting the enriched data in a structured data warehouse (Google BigQuery), validated by 11 automated integrity tests.

The end result is a centralized, analysis-ready 'single source of truth' that enables iO to generate in-depth benchmarks, portfolio-wide trends, and sector-specific insights. The estimated operational costs of the platform, at €30-€200 per month, are extremely efficient in relation to its strategic value.

The concrete technical implementation of this plan is documented in the Implementation chapter.

Table of Contents

<i>Project Proposal and Methodological Justification</i>	<i>3</i>
<i>Management Summary</i>	<i>3</i>
<i>1. Introduction and Problem Context</i>	<i>6</i>
1.1. Central Research Question	6
<i>2. Context Analysis: Organization and Technological Domain</i>	<i>7</i>
2.1. Analysis of the Organizational Context: iO	7
2.2. Analysis of the Technological Domain: Google Analytics 4	7
2.3. Strategic Facilitators: Partner Statuses	8
<i>3. Problem Statement</i>	<i>9</i>
3.1. Current Situation: Data Richness in Silos	9
3.2. The Core Problem and its Consequences	9
3.3. Critical Time Factor: The Risk of Permanent Data Loss	9
<i>4. Objectives and Scope.....</i>	<i>10</i>
4.1. Main Objective	10
4.2. Sub-objectives	10
4.3. Scope Definition	10
4.3.1. In Scope.....	11
4.3.2. Out of Scope	11
<i>5. Strategic Value and Return on Investment.....</i>	<i>12</i>
5.1. Value Proposition per Stakeholder	12
5.1.1. Strategic Management	12
5.1.2. Marketing and Data Consultants	12
5.1.3. Clients	12
5.2. Return on Investment (ROI) Analysis	13
<i>6. Preliminary Research and Architectural Decision-Making.....</i>	<i>14</i>
6.1. Data Volume Analysis and Projection	14
6.2. Comparative Analysis of Database Technologies	14
6.3. Economic Analysis: Debunking the Myth of High Costs	15
6.4. Derived Architectural Design	15
<i>7. Architectural Design.....</i>	<i>17</i>
7.1. Architectural Principles	17
7.2. Conceptual Model: Layered Data Processing.....	18

7.3.	Design for Non-Functional Requirements	18
8.	<i>Project Planning and Phasing</i>	19
8.1.	Overzicht	19
8.2.	Phasing and Deliverables.....	20
	Phase 1 — Research (Weeks 1–2).....	20
	Phase 2 — Pipeline (Week 3)	20
	Phase 3 — Enrichment (Week 4).....	20
	Phase 4 — Analysis (Week 5).....	20
	Phase 5 — BigQuery (Week 6)	21
	Phase 6 — Bonzai Integration (Weeks 7–8).....	21
9.	<i>Communication and Reporting Plan</i>	22
9.1.	Stakeholders and Information Needs	22
9.2.	Communication Matrix.....	22
9.3.	Layered Traceability Strategy	23
10.	<i>Risk Analysis and Mitigation Strategies</i>	25
10.1.	Technical Risks.....	25
10.2.	Operational Risks	25
10.3.	Organizational and Legal Risks.....	26
11.	<i>Evaluation Framework and Success Criteria</i>	27
11.1.	Evaluation Matrix.....	27
12.	<i>References</i>	29
12.1.	Data Sources and APIs	29
12.2.	Database Technologies	29
12.3.	Methodology and Articles	29

1. Introduction and Problem Context

In the current digital marketing landscape, the ability to make data-driven decisions is a critical success factor. Organizations that provide large-scale digital services, such as iO, manage an immense and continuously growing amount of data about user interactions. For iO, with a portfolio of hundreds of clients in Northern Europe, this data primarily manifests within the Google Analytics 4 (GA4) ecosystem. This ecosystem represents a strategic, yet largely untapped, source of information.

The current practice for data analysis within iO is fragmented. Insights are generated per individual client, which makes cross-client analyses and portfolio-wide benchmarking impossible. Furthermore, the limited data retention of the GA4 platform (maximum of 14 months) poses a significant risk of permanent data loss. These factors lead to a suboptimal utilization of available data and a missed opportunity for strategic value creation.

This bachelor's project addresses this challenge through the conception, design, and implementation of a scalable, automated data pipeline. The core function of this system is the systematic extraction, transformation, and loading (ETL) of GA4 data into a centralized data warehouse. By enriching the raw data with crucial contextual dimensions – such as an AI-driven sector classification – an analysis-ready 'single source of truth' is created.

1.1. Central Research Question

The central research question that this project aims to answer is as follows:

How can a scalable and robust data pipeline be designed and implemented that enables the systematic extraction, enrichment, and centralization of Google Analytics 4 data from a diverse and extensive client portfolio, in order to facilitate portfolio-wide analyses and strategic insights?

2. Context Analysis: Organization and Technological Domain

A thorough understanding of both the organizational context and the technological domain is essential for correctly positioning this project. This analysis covers the two pillars that form the basis of the problem statement: the organization iO and the Google Analytics 4 ecosystem.

2.1. Analysis of the Organizational Context: iO

iO operates as a leading, full-service digital agency with a strong presence in Northern Europe. By delivering a broad spectrum of services – from strategy and technology to marketing and creativity – iO serves a diverse client portfolio, ranging from small and medium-sized enterprises to multinationals.

This business structure has led to the management of an extensive and heterogeneous landscape of digital assets. Specifically, this translates to:

- Hundreds of Google Analytics accounts.
- Thousands of unique properties (websites, webshops, mobile applications).
- A client base spread across various sectors (e.g., e-commerce, healthcare, automotive, financial services).

This scale creates a unique, strategic data asset. The aggregated data from this portfolio represents a cross-section of the digital economy that would be impossible for an individual client to reconstruct. Unlocking this asset for cross-client benchmarking and portfolio-wide trend analysis forms the core of the opportunity this project addresses.

2.2. Analysis of the Technological Domain: Google Analytics 4

Google Analytics 4 (GA4) has been Google's dominant analytics platform since 2023 and is the primary data source for this project. The relevance of GA4 is determined by several fundamental characteristics:

- **Event-based Data Model:** Unlike its predecessors, GA4 is fundamentally event-based. Every user interaction, from a page visit to a purchase, is registered as a uniform 'event'. This model provides a significantly richer and more granular dataset but also requires a new approach to data aggregation and analysis.
- **Cross-platform Data Collection:** GA4 is designed to unify data from both websites and mobile applications within a single property, enabling a holistic view of the user.
- **Natural Ecosystem Integration:** As a Google product, GA4 offers seamless, native integration with other crucial components of the Google Cloud Platform, most notably the data warehouse Google BigQuery. This integration is a key factor that facilitates the technical feasibility of a scalable data export.

2.3. Strategic Facilitators: Partner Statuses

The technical and strategic feasibility of the project is further enhanced by iO's formal partner statuses:

Partner	Status	Strategic Relevance
Google	Premium Partner	This status offers direct benefits such as access to Google Cloud Platform (GCP) credits, which significantly reduces initial development and testing costs. It also provides access to specialized engineering support, which can serve as a valuable knowledge resource.
AWS	Advanced Partner	Although the project primarily relies on the Google ecosystem, this status offers strategic flexibility for future expansions, for example, for application hosting or the use of alternative machine learning models via Amazon Bedrock.

These factors combined – a unique data asset within the organization, a powerful technological platform, and strategic partnerships – create the ideal foundation for the successful realization of this project.

3. Problem Statement

The context analysis shows that iO possesses an extensive and strategically valuable data asset. However, current methods for data analysis and management are unable to unlock this value. This chapter articulates the core problem that arises from this situation.

3.1. Current Situation: Data Richness in Silos

The current situation is characterized by a paradox: an abundance of data, but a scarcity of insight. Each client's data is isolated within its own Google Analytics 4 interface. Access is manual, fragmented, and not scalable. These data silos form a fundamental barrier to any form of aggregated, portfolio-wide analysis.

3.2. The Core Problem and its Consequences

The central problem is the lack of an automated, centralized system for data extraction, enrichment, and storage. This technical shortcoming leads to a series of direct, negative business consequences:

1. **Inability for Cross-Client Benchmarking:** It is impossible to benchmark a client's performance against that of industry peers. Questions like "How does client X's conversion rate perform compared to the sector average?" cannot be answered. This limits the strategic advisory value that iO can offer.
2. **Lack of Portfolio-wide Insights:** Subtle trends, patterns, or anomalies that are only visible at a macro level (for example, a sudden drop in traffic within a specific sector) go unnoticed. This leads to missed opportunities for proactive intervention and strategic planning.
3. **Reactive instead of Proactive Service Delivery:** Without automated monitoring and central data analysis, consultants can often only react after a problem (e.g., a technical issue affecting tracking) is reported by the client.
4. **Underutilization of a Strategic Data Asset:** iO's unique, aggregated dataset, which represents a cross-section of the digital economy, remains completely untapped. The potential value for internal knowledge development, product innovation, and thought leadership is not being realized.

3.3. Critical Time Factor: The Risk of Permanent Data Loss

The urgency of this problem is significantly amplified by a technical limitation of the Google Analytics 4 platform: the maximum data retention of 14 months. All data older than this period is irrevocably deleted by Google. Every day without a structural solution for data extraction results in the permanent loss of valuable historical data. This creates a compelling need to implement a sustainable data archiving solution in the short term.

4. Objectives and Scope

In response to the formulated problem statement, this chapter defines the concrete and measurable objectives of the project. The project scope is also precisely defined.

4.1. Main Objective

The main objective of this project is:

To design, implement, and validate a scalable and automated data pipeline that transforms Google Analytics 4 (GA4) data from the entire iO client portfolio into a centralized, enriched, and analysis-ready data asset.

4.2. Sub-objectives

To achieve this main objective, the following four concrete and measurable sub-objectives have been formulated:

1. Realization of an Automated ETL Process:

Developing a robust Python application that automates the entire ETL (Extract, Transform, Load) process for GA4 data. This includes data extraction via the official APIs, the application of an intelligent filtering system to ensure data quality, and the structured storage of the raw data in a local database.

2. Implementation of Advanced Data Enrichment:

Increasing the analytical value of the data by implementing various enrichment mechanisms. The most important of these are (a) an AI-driven classification that assigns each account to one of 25 defined sectors, and (b) the addition of more than 14 contextual dimensions, including geography, device category, acquisition channel, and revenue data.

3. Validation of a Scalable Cloud Architecture:

Demonstrating that the developed solution is technically feasible and scalable in a production environment. This is validated by a successful pilot migration of a representative dataset to Google BigQuery, including the implementation of performance optimizations (partitioning, clustering) and the successful execution of an automated validation protocol with 11 unique checks.

4. Delivery of Initial Analysis Artifacts and Documentation:

Unlocking the value of the created data asset by delivering a set of Jupyter Notebooks with initial analyses and visualizations. In addition, complete technical and project documentation will be delivered to ensure the transferability and maintenance of the solution.

4.3. Scope Definition

A clear definition of the scope is crucial for successfully completing the project within the set deadline.

4.3.1. In Scope

The project includes the complete end-to-end realization of the data pipeline, from authentication with the Google APIs to the validated storage in BigQuery. This includes the development of ETL scripts, the AI-driven sector classification, data enrichment with more than 14 dimensions, and the implementation of a robust checkpoint system. The validation is supported by an automated test protocol. Finally, the scope includes the creation of initial analysis notebooks and a thorough cost analysis for the production environment.

4.3.2. Out of Scope

To ensure focus and feasibility, the following elements are explicitly placed out of scope:

- **Real-time Data Processing:** A batch-oriented architecture (e.g., daily) is deliberately chosen to manage technical complexity and operational costs. Real-time streaming has been identified as a possible, but separate, future extension.
- **Adjustments in GA4 Configurations:** The project consumes the data as it is provided. Advising clients on or implementing changes to their tracking setup is a consultancy activity and falls outside this technical project.
- **Development of BI Dashboards:** Although the delivered data asset is the ideal source for dashboarding (e.g., in Looker Studio), the actual creation of these dashboards is out of scope. The delivery of the validated dataset is the endpoint of this project.
- **Predictive Analytics:** The focus is on unlocking and descriptively analyzing the data. The development of predictive models is considered a potential next step but is not part of this project.

5. Strategic Value and Return on Investment

The realization of the GA4 Data Pipeline transcends a purely technical implementation; the project delivers direct and measurable strategic value for the entire organization. This chapter analyzes the value proposition for the primary stakeholders and substantiates the Return on Investment (ROI).

5.1. Value Proposition per Stakeholder

The created data asset generates specific value for three crucial groups within and outside the organization:

5.1.1. Strategic Management

For iO's management, the pipeline transforms data from a fragmented, operational source into a strategic management tool. It provides a macroeconomic bird's-eye view of the entire client portfolio, bringing previously unanswerable questions within reach:

- **Portfolio Management:** Which sectors are showing the most growth? Which client segments are lagging behind?
- **Resource Allocation:** Where is the greatest optimization potential, and how can we align our resources (people, budget) accordingly?
- **Strategic Decision-Making:** Decisions about market trends, investments, and service development can be substantiated with quantitative data instead of relying solely on intuition.

5.1.2. Marketing and Data Consultants

For the consultants who work with clients daily, the pipeline delivers a direct improvement in efficiency and quality. Their role evolves from reactive to proactive.

- **Increased Advisory Value:** Reporting is enriched with context. A conversion rate of 2.3% is no longer an abstract number but can be interpreted as "significantly above the sector average of 1.8%."
- **Efficiency Improvement:** The time previously spent on manual data export and analysis can now be invested in delivering high-value strategic advice.
- **Proactive Signaling:** The central dataset makes it possible to detect anomalies (e.g., a sudden drop in mobile traffic for an entire sector) before the client notices, creating opportunities for proactive service delivery and upselling.

5.1.3. Clients

Although iO's clients are not direct users of the pipeline, they are the ultimate beneficiaries of the improved service delivery. They receive more advanced, data-substantiated reports that not only describe their own performance but also place it in

the context of their market. This increases the perceived value of iO's services and strengthens the client relationship.

5.2. Return on Investment (ROI) Analysis

The return on investment is analyzed by comparing the relatively low operational costs with the significant, quantified, and qualified benefits.

- **Operational Costs:** The estimated monthly costs for the platform, including data storage in BigQuery, hosting, and API usage, are approximately €30 to €200. This amounts to an annual cost of only €360 to €2,400.
- **Quantified Benefits (Efficiency Gains):** A significant, directly measurable benefit is the time savings for consultants. A conservative estimate of 30 minutes saved per monthly client report, extrapolated over 200 clients, results in an annual efficiency gain of approximately 1,200 consultancy hours. These hours can be reinvested in billable, value-adding activities.
- **Qualified Benefits (Strategic Impact):** The most substantial value lies in the strategic impact. The improved service delivery and proactive insights increase client satisfaction and loyalty. If the implementation of this platform contributes to retaining just one to two medium-sized clients (with contract values often between €50,000 and €200,000 per year), the return exceeds the investment by a factor of 20 to 100. Furthermore, it positions iO as a data-driven thought leader in the market, which has a positive impact on new client acquisition.

6. Preliminary Research and Architectural Decision-Making

The choice of a specific technological architecture is one of the most critical decisions in a data engineering project. This choice must not only be based on technical specifications but also on a thorough analysis of data volumes, costs, and the specific requirements of the use case. This chapter documents the preliminary research and the resulting decision-making that underlies the chosen architecture.

6.1. Data Volume Analysis and Projection

An accurate estimation of the expected data volume is fundamental to any architectural decision. To this end, a methodological projection was carried out.

Methodology: A Python script (`estimate_database_size.py`) analyzes a representative sample of accounts and classifies them into four data volume categories (Small, Medium, Large, Mega). To form a realistic picture, a weighted extrapolation was applied to the entire portfolio of x accounts. This weighting corrects for the inevitable selection bias in the initial testing phase, in which data-rich accounts were overrepresented. A realistic portfolio distribution was assumed (40% small, 35% medium, 18% large, 7% mega), which is representative of a typical digital agency.

Results: The analysis resulted in three scenarios:

Scenario	Monthly Rows (enriched)	Estimated BigQuery Volume (30d)
Conservative	~87 million	~13 GB
Realistic	~143 million	~21 GB
Worst-case	~300 million	~45 GB

Conclusion: The expected monthly data influx is substantial, with a realistic estimate of 21 GB. With a 14-month retention, the dataset grows to approximately 294 GB. This volume is large enough to make manual analysis impossible, but—and this is a crucial finding—it falls comfortably within the 'on-demand' pricing models of modern cloud data warehouses.

6.2. Comparative Analysis of Database Technologies

With the volume projection as a starting point, five database technologies were evaluated based on a set of weighted criteria, including cost, performance, scalability, operational management, and the seamlessness of integration with the existing Google ecosystem (GA4, Looker Studio).

- Google BigQuery (On-demand): Evaluated as the superior choice for the production database. Its serverless architecture, petabyte scalability, and 'pay-as-you-go' pricing model (€1-10/month for our volume) are decisive. The native

integration with GA4 and Looker Studio significantly reduces technical complexity and accelerates time-to-market.

- **DuckDB (Embedded OLAP):** Selected as the query engine for the Bonzai AI integration. DuckDB's ability to read SQLite files directly in-process, combined with sub-100ms query times for analytical queries, makes it the ideal, cost-free solution for fast, synchronous queries from an application backend. However, it is not a suitable replacement for BigQuery due to the lack of a native Looker Studio connector and its limitations for multi-user access.
- **SQLite:** Retains its role as the development and staging database. It is ideal for local development and functions as a reliable intermediate station for the ETL pipeline, but it lacks the scalability and concurrent user capacity required in a production environment.
- **Databricks & AWS Redshift:** Both were rejected as disproportionate. Their minimum monthly costs (€300-€720+) and the architectural overhead for setting up a data ingestion pipeline from the Google ecosystem are not in proportion to the current data volume. These platforms are designed for terabyte or petabyte scale and would be a form of 'over-engineering' at this stage.

6.3. Economic Analysis: Debunking the Myth of High Costs

A significant result of the preliminary research was the debunking of the initial assumption that the project would entail high query costs. This assumption was based on the size of the raw data export that GA4 can natively produce (1-3 TB).

However, the research showed that by using the GA4 Data API, which returns aggregated data, the resulting dataset is 50 to 100 times smaller. This reduces the cost of a full table scan from a potential €18.75 to just €0.13 - €0.28. This crucial finding makes BigQuery's 'on-demand' pricing model economically extremely viable.

The total estimated platform costs, including BigQuery, application hosting, and API usage, are projected to be €30 to €200 per month. This is a fraction of the cost of alternative, fixed-capacity-based models (such as BigQuery Editions or Databricks), which would be more than 200 times too expensive for this project.

6.4. Derived Architectural Design

Based on the preceding analyses, a hybrid, 'best-of-breed' architecture has been designed that maximizes cost-efficiency and performance:

1. **ETL Pipeline (Python):** The existing, validated Python application remains the heart of the system, responsible for extraction, enrichment, and staging in a local SQLite database.
2. **Persistence & Dashboarding (BigQuery):** The enriched data is periodically migrated to Google BigQuery. BigQuery serves as the persistent 'single source of


truth' and as the data source for BI tools like Looker Studio, for which it offers a native, high-performance connector.

3. Real-time AI Queries (DuckDB): For the Bonzai AI integration, which requires direct answers to natural language questions, DuckDB is used. It reads the most recent SQLite data in-memory and delivers lightning-fast query results without network latency or additional costs.

This architecture is based on the principle of proportionality: use the right tool for the right job, and avoid unnecessary costs and complexity. It maximizes the benefits of the Google ecosystem (GA4 → BigQuery → Looker Studio) while simultaneously leveraging the speed and efficiency of an embedded engine like DuckDB where it makes the most sense.

7. Architectural Design

Building upon the preliminary research and the technological choices made, this chapter describes the conceptual and logical design of the data pipeline. This architectural design forms the blueprint for the technical realization of the system.



7.1. Architectural Principles

The design is based on several fundamental principles that ensure the quality, sustainability, and efficiency of the solution:

- **Modularity & Separation of Concerns:** The system is divided into logical, independent components (e.g., authentication, data extraction, data enrichment), which increases maintainability and testability.
- **Proportional Costs:** The architecture is designed to keep operational costs proportional to data volume and usage by deliberately choosing serverless and 'on-demand' technologies.
- **Maximum Automation:** All steps in the ETL process, from data extraction to validation, are designed to function without manual intervention.
- **Fault Tolerance and Robustness:** The system is designed to operate under unpredictable conditions, such as API failures or network interruptions.

7.2. Conceptual Model: Layered Data Processing

The architecture can be modeled as a layered system in which data is sequentially processed and enriched. Each layer has a specific responsibility in the ETL process.

- **Layer 1: Data Ingestion:** This layer is responsible for communication with external APIs. It discovers the digital assets (properties) to be processed and collects the raw, aggregated event data.
- **Layer 2: Data Processing & Enrichment:** In this core layer, the raw data is transformed into a valuable dataset. This includes filtering out noise, assigning a sector via AI, and adding more than 14 analytical dimensions.
- **Layer 3: Data Persistence:** The processed data is stored. SQLite serves as a lightweight, reliable staging environment. Google BigQuery acts as the scalable, permanent data warehouse for production and BI purposes. DuckDB is used as a high-performance in-memory query engine for specific, real-time application needs (Bonzai).

7.3. Design for Non-Functional Requirements

The design explicitly addresses a number of critical non-functional requirements:

- **Fault Tolerance:** The pipeline is 'stateful' thanks to a checkpoint system. In case of an interruption, the process can be resumed without data loss or duplicate processing. API communication is wrapped in an 'exponential backoff' mechanism to handle temporary failures.
- **Data Integrity:** All write operations to the database are designed to be idempotent (using INSERT OR IGNORE), which guarantees that repeated executions of the same task do not lead to data duplication.
- **Scalability:** The architecture is horizontally scalable. By choosing serverless components like BigQuery and batch processing per property, the system can grow with an increasing number of clients without fundamental redesigns.

The concrete implementation of this design, including the detailed pipeline steps and the database schema, is described in the Implementation chapter.

8. Project Planning and Phasing

The project was executed according to a structured, phased model. Each phase builds on the results of the previous one and delivers a concrete set of deliverables.

8.1. Overzicht

#	Period	Phase	Focus	Key Result	Status
1	Week 1–2	Research	Getting to know GA4, GTM, and related tools	Working knowledge of the GA4 ecosystem and tooling	Completed
2	Week 3	Pipeline	Retrieving and storing data	Working data pipeline with SQLite storage	Completed
3	Week 4	Enrichment	Adding context to the data	Enriched dataset with sector, device, geo, channel, and revenue	Completed
4	Week 5	Analysis	Extracting insights from the enriched data	8 analysis notebooks with visualizations and benchmarks	Completed
5	Week 6	BigQuery	Migration from local SQLite to a scalable cloud database	Validated BigQuery migration with cost analysis	Active
6	Week 7–8	Bonzai Integration	Embedding the pipeline into the existing Bonzai platform	Live integration within iO's Bonzai platform	Active

8.2. Phasing and Deliverables

Phase 1 — Research (Weeks 1–2)

The goal of this initial phase was to lay a solid theoretical and technical foundation. This included an in-depth study of the GA4 ecosystem and associated technologies, and the validation of the initial technical feasibility.

- Deliverable: A Proof of Concept demonstrating successful data extraction via the Google APIs.
- Deliverable: A complete inventory of the accounts and properties to be processed.

Phase 2 — Pipeline (Week 3)

In this phase, the foundation of the ETL pipeline was developed, focusing on robust and scalable extraction and storage of raw data.

- Deliverable: A functional ETL pipeline that automatically performs the Discover and Fetch steps.
- Deliverable: A normalized SQLite database that serves as a staging area.
- Deliverable: Implementation of robustness mechanisms (e.g., exponential backoff and a stateful checkpoint system).

Phase 3 — Enrichment (Week 4)

This phase focused on transforming raw data into valuable information by adding contextual dimensions.

- Deliverable: Integration of an AI module for the automated classification of accounts into 25 sectors.
- Deliverable: An enrichment module that expands the dataset with more than 14 dimensions (e.g., geography, device, acquisition channel, revenue).

Phase 4 — Analysis (Week 5)

The goal of this phase was to unlock the value of the created dataset through concrete analyses.

- Deliverable: Delivery of 8 Jupyter Notebooks with analyses ranging from operational monitoring to strategic sector benchmarks.
- Deliverable: An export module that makes the enriched data available in CSV format.

Phase 5 — BigQuery (Week 6)

This phase focused on validating the scalability and production viability of the solution in a cloud environment.

- Deliverable: A successful pilot migration of a representative dataset to Google BigQuery.
- Deliverable: An automated validation protocol with 11 unique checks that ensures data integrity between source and target systems.
- Deliverable: A detailed cost analysis with substantiated recommendations for the production infrastructure.

Phase 6 — Bonzai Integration (Weeks 7–8)

The final phase focuses on integrating the pipeline into existing business processes and completing the project documentation.

- Planned Deliverable: Integration with the internal Bonzai platform to make insights available to consultants in real-time.
- Planned Deliverable: Completion of the technical documentation and user manual.
- Planned Deliverable: Preparation and execution of the final presentation and demonstration.

9. Communication and Reporting Plan

A structured approach to communication and reporting is essential for the success of any project. This plan defines the methods and frequencies for monitoring progress and informing stakeholders, and describes the strategy for ensuring the traceability of both project-related and technical activities.

9.1. Stakeholders and Information Needs

The communication strategy is tailored to the information needs of the primary stakeholders:

- Project Supervisor (iO): Requires detailed insight into daily progress, technical challenges, and the planning of next steps to provide timely and effective guidance.
- Academic Supervisor (University): Requires periodic, formal reports on progress, achieved learning objectives, and the alignment of the project with the academic final terms.
- Project Team / Management (iO): Requires periodic updates on achieved milestones and the overall status of the project to monitor alignment with business objectives.

9.2. Communication Matrix

The following matrix formalizes the communication flows within the project:

Communication Tool	Stakeholder(s)	Frequency	Purpose
Daily Stand-up	Project Supervisor	Daily	Discussing progress, identifying blockers, asking ad-hoc questions.
Weekly Meeting	Project Supervisor	Weekly	Evaluating the past week, planning the upcoming week, and in-depth feedback sessions.
Progress Report (E-mail)	Academic Supervisor	Bi-weekly	Formal, written reporting of project progress, in accordance with

			the educational institution's requirements.
Milestone Updates (Slack/E-mail)	Project Team / Management	Upon phase completion	Informing the broader team about the successful completion of a project phase and the achieved results.
Jira (Project Management Tool)	Project Supervisor	Continuous	Asynchronous management of tasks, specifications, and the status of individual work packages.

9.3. Layered Traceability Strategy

To ensure full transparency and reproducibility, the project uses a layered traceability strategy that covers both project progress and technical execution.

- **Project Traceability:** Is ensured by a combination of:
 - Jira Tickets: Each work package, from research task to bug fix, is defined as a ticket with a clear status and scope.
 - GitHub Commits: Every code change is linked to a clear commit message, which often refers to a specific Jira ticket.
 - Bridge Time Tracking: The daily registration of hours worked and activities provides a detailed log of time spent.
- **Technical Traceability:** Is ensured by the automated logging of the application itself:
 - Application Logs: The pipeline generates structured log files with timestamps, runtime statistics, and any error messages (including stack traces).
 - Database Checkpoints: The pipeline_runs table in the database functions as a permanent, auditable log of each pipeline execution, including the status per processed property. This is essential for both debugging and the 'stateful recovery' mechanism.

This combined approach ensures that every activity within the project, from a decision made to a line of code executed, is fully traceable and accountable.

10. Risk Analysis and Mitigation Strategies

Proactive risk management is a crucial part of successful project management. This chapter identifies and categorizes the main risks for the project and defines the mitigation strategies designed to minimize the impact and likelihood of these risks. Many of these strategies have already been integrated as preventive measures in the system architecture.

10.1. Technical Risks

These are risks inherent to the technology and the external systems with which it integrates.

- Risk: Exceeding API Quotas.
 - Analysis: The Google Analytics APIs have strict quotas. During a full production run, there is a chance that these limits will be reached, which could interrupt the pipeline.
 - Mitigation Strategy: A multi-layered strategy has been implemented. (1) Proactive Rate Limiting: The pipeline pauses briefly between API calls. (2) Reactive Error Handling: An 'exponential backoff' mechanism with jitter ensures an intelligent retry attempt for quota-related errors. (3) Stateful Recovery: The checkpoint system ensures that the pipeline can be resumed after a hard stop without losing work already performed.
- Risk: Unexpected API Changes.
 - Analysis: External APIs can change (breaking changes), which can disrupt the functionality of the pipeline.
 - Mitigation Strategy: The application uses specified API versions. The code that interacts with the APIs is centralized in specific client modules, allowing adjustments to be made in an isolated and efficient manner.

10.2. Operational Risks

These are risks related to the execution and management of the system at production scale.

- Risk: Insufficient Pipeline Performance at Scale.
 - Analysis: The processing time for the entire portfolio of x accounts could exceed the daily batch window.
 - Mitigation Strategy: The architecture is prepared for parallelization, allowing multiple accounts to be processed simultaneously. In addition, the data retrieval is designed to be incremental, so that after the initial run, only the most recent data needs to be fetched, drastically reducing daily execution time.

- Risk: Uncontrolled Data Growth (Dimension Explosion).
 - Analysis: Fields with very high cardinality (e.g., landingPage, city) can cause the database size to grow unexpectedly fast.
 - Mitigation Strategy: Data size is actively monitored with a specific utility script (estimate_database_size.py). If necessary, high-cardinality fields can be excluded from enrichment via the central configuration file without requiring a code change.
- Risk: Inconsistency in AI Classification.
 - Mitigation Strategy: This risk is minimized by (1) setting the 'temperature' parameter of the AI model to 0.0 for maximum determinism, (2) forcing the model to choose from a fixed list of 25 sectors, and (3) implementing a manual override.csv file that acts as the 'final word'.

10.3. Organizational and Legal Risks

- Risk: Complexity of Platform Integration.
 - Analysis: The integration of the pipeline with the existing, complex Bonzai platform may take longer than planned.
 - Mitigation Strategy: A phased 'Minimum Viable Product' (MVP) approach has been chosen. The initial integration focuses on a simple, embedded DuckDB solution. The more complex, production-ready connection to the BigQuery backend will be developed in a later phase.
- Risk: Privacy and GDPR Implications.
 - Analysis: Centralizing data from different clients in a single database can raise questions regarding data protection.
 - Mitigation Strategy: The pipeline's design proactively mitigates this risk. No Personally Identifiable Information (PII) is processed; all data is aggregated at the level of event counts. For benchmarking purposes, 'Gold tables' are created that further aggregate data to the sector level, completely abstracting individual client data.

11. Evaluation Framework and Success Criteria

The success of this bachelor's project is not assessed subjectively but is tested against a set of predefined, measurable success criteria. This evaluation framework directly links each criterion to one of the project objectives as formulated in Chapter 4, and defines how the achievement of each criterion is objectively verified.

11.1. Evaluation Matrix

Project Objective (ref. Ch. 4)	Success Criterion	Verification Method	Status
1. Realization of Automated ETL Process	The pipeline can process the entire portfolio end-to-end without data loss during interruptions.	A full run over all accounts is successfully completed, with the checkpoint system functioning correctly after a simulated interruption.	✓ Achieved
1. Realization of Automated ETL Process	The system is idempotent; repeated runs do not lead to data duplication.	Executing a second, identical pipeline run results in 0 new rows in the database, verified via an SQL query.	✓ Achieved
1. Realization of Automated ETL Process	The pipeline manages resources efficiently by ignoring inactive properties.	Properties that do not generate data for N runs automatically get the status <code>is_active = 0</code> in the properties table.	✓ Achieved
2. Implementation of Advanced Data Enrichment	All active properties are programmatically classified into one of the 25 defined sectors.	An SQL query (<code>SELECT COUNT(*) FROM properties WHERE sector IS NOT NULL</code>) shows that all active properties have a sector label.	✓ Achieved

2. Implementation of Advanced Data Enrichment	The raw event data is enriched with at least 14 additional dimensions.	The ga4_events_enriched table contains the specified columns for geography, device, acquisition channel, and revenue.	✓ Achieved
3. Validation of Scalable Cloud Architecture	The data model and ETL logic are successfully migratable and scalable in a production environment.	The pilot migration to Google BigQuery, including the automated validation protocol with 11 tests, succeeds without errors.	✓ Achieved
4. Delivery of Analysis Artifacts & Documentation	The strategic value of the data is demonstrated through concrete analyses.	Delivery of 8 Jupyter Notebooks that generate portfolio-wide insights and visualizations based on the created data asset.	✓ Achieved
4. Delivery of Analysis Artifacts & Documentation	The architectural choices are substantiated by thorough preliminary research.	The project documentation includes a formal comparison of database technologies and a detailed cost analysis (Chapter 6).	✓ Achieved
4. Delivery of Analysis Artifacts & Documentation	The project result is transferable and integrable into the organization.	Delivery of a Minimum Viable Product (MVP) for the Bonzai integration, which can perform live queries on the dataset.	🔄 In progress

12. References

This chapter provides an overview of the main external sources, documentation, and articles consulted during the analysis and design phase of the project.

12.1. Data Sources and APIs

- Google Analytics Data API: Official Documentation
- Google Analytics Admin API: Official Documentation

12.2. Database Technologies

- Google BigQuery: Product Overview and Documentation
- DuckDB: Official Documentation
- SQLite: Official Website

12.3. Methodology and Articles

- Google Cloud, "BigQuery pricing": Pricing Model Documentation
- DuckDB, "Modern Data Stack in a Box": Blog Post